

# Importing and installing required libraries

```
In [18]: import pandas as pd
from datetime import datetime
from matplotlib import pyplot
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

## Question 1

(8 pts) The Bureau of Transportation Statistics (BTS) conducted a study to evaluate the impact of Sept 11 attacks (9/11) on U.S. air transportation. The purpose of this study is to provide a greater understanding of the passenger travel behavior patterns of persons travelling by air before and after the event. In order to assess the impact of September 11, BTS took the following approach: Using data before September 11, it forecasted future data (under the assumption of no terrorist attack). Then, BTS compared the forecasted series with the actual data to assess the impact of the event.

## Importing, examining and cleaning the BTS Dataset

In [2]: *# importing the dataset*

```
df_bts = pd.read_csv("BTS_Air_Rail_Vehicle_Miles.csv")
df_bts
```

Out[2]:

	Month	Air	Rail	Vehicle
0	Jan-90	35153577	454115779	163.28
1	Feb-90	32965187	435086002	153.25
2	Mar-90	39993913	568289732	178.42
3	Apr-90	37981886	568101697	178.68
4	May-90	38419672	539628385	188.88
...	...	...	...	...
167	Dec-03	57795908	489403554	237.60
168	Jan-04	53447972	410338691	217.30
169	Feb-04	52608801	389778365	210.40
170	Mar-04	63600019	453014590	247.50
171	Apr-04	61887720	471116666	245.40

172 rows × 4 columns

In [3]: *# data types*

```
df_bts.dtypes
```

```
Out[3]: Month      object
Air          int64
Rail         int64
Vehicle     float64
dtype: object
```

```
In [4]: df_bts.dtypes
```

```
Out[4]: Month      object  
        Air        int64  
        Rail       int64  
        Vehicle   float64  
        dtype: object
```

```
In [5]: df_bts.columns
```

```
Out[5]: Index(['Month', 'Air ', 'Rail', 'Vehicle'], dtype='object')
```

## Part a

Is the goal of this study descriptive or predictive?

Although, there is some aspect of the study involving forecasting, the ultimate goal is descriptive analysis. The problem involves forecasting travel behaviours post 9/11 first. Then these values are compared to the ground truth values. This involves descriptive analysis. Hence, the goal of the study is largely descriptive.

## Part b

Create a time series plot of the Air data, i.e. a plot  $y_t$  versus  $t$ , where  $t=1,2,3 \dots$ . What would  $t=1, 2, 3$  refer to in the time series? Which time period does  $t=1$  refer to?


```
In [6]: df_bts_air = df_bts[['Month', 'Air ']]  
df_bts_air.set_index('Month', inplace=True)  
df_bts_air
```

Out[6]:

Air	
Month	
Jan-90	35153577
Feb-90	32965187
Mar-90	39993913
Apr-90	37981886
May-90	38419672
...	...
Dec-03	57795908
Jan-04	53447972
Feb-04	52608801
Mar-04	63600019
Apr-04	61887720

172 rows × 1 columns

```
In [7]: plt.rcParams["figure.figsize"] = (15,5)
df_bts_air.plot()
plt.xticks(rotation = 45)
plt.show()
```



$t = 1, 2, 3, \dots$  represent to the month, year in the time series.  $t = 1$  refers to the month of January in the year 1990.

## Part c

What are the values for  $y_1$ ,  $y_2$  and  $y_3$  in the time series?

$y_1 = 35153577$

$y_2 = 32965187$

$y_3 = 39993913$

## Question 2


(10 pts) In addition to air travel data, two additional time series are also provided in the same data file – Rail and Vehicle travel.

### Part a

Which of these components appear in the Air and Vehicle time series: i) Level; ii) Seasonality; iii) Trend; iv) Noise. List for each data set.

### Components of the air time series

```
In [8]: df_bts_air = df_bts[['Month', 'Air ']]
df_bts_air.set_index('Month', inplace=True)
plt.rcParams["figure.figsize"] = (15,5)
df_bts_air.plot()
plt.xticks(rotation = 45)
plt.show()
```




Components appearing in the air time series:

1. Level
2. Seasonality
3. Trend
4. Noise

## Components of the vehicle time series

```
In [10]: df_bts_vehicle = df_bts[['Month', 'Vehicle']]
df_bts_vehicle.set_index('Month', inplace=True)
plt.rcParams["figure.figsize"] = (15,5)
df_bts_vehicle.plot()
plt.xticks(rotation = 45)
plt.show()
```






Components appearing in the vehicle time series:

1. Level
2. Seasonality
3. Trend

## Part b

For the Rail data set, describe the trend, i.e. how does the trend vary across the time series?

```
In [42]: df_bts_air = df_bts[['Month', 'Rail']]
df_bts_air.set_index('Month', inplace=True)
plt.rcParams["figure.figsize"] = (15,5)
df_bts_air.plot()
plt.xticks(rotation = 45)
plt.show()
```



In the rail dataset, the trend does not remain constant across the entire time series. There is a clearly observable downward trend until

1998-05. After that, there is a noticeable upward trend.

## Question 3

(6 pts) Forecasting Shampoo Sales: The file ShampooSales.csv contains data on the monthly sales of a certain shampoo over a 3 year period.

### Importing, examining and cleaning the shampoo sales dataset

In [46]: *# Importing the dataset*

```
df_shampoosales = pd.read_csv('ShampooSales.csv')  
df_shampoosales
```

Out[46]:

	Month	Shampoo Sales
0	Jan-95	266.0
1	Feb-95	145.9
2	Mar-95	183.1
3	Apr-95	119.3
4	May-95	180.3
5	Jun-95	168.5
6	Jul-95	231.8
7	Aug-95	224.5
8	Sep-95	192.8
9	Oct-95	122.9
10	Nov-95	336.5
11	Dec-95	185.9
12	Jan-96	194.3
13	Feb-96	149.5
14	Mar-96	210.1
15	Apr-96	273.3
16	May-96	191.4
17	Jun-96	287.0
18	Jul-96	226.0
19	Aug-96	303.6
20	Sep-96	289.9
21	Oct-96	421.6

	Month	Shampoo Sales
22	Nov-96	264.5
23	Dec-96	342.3
24	Jan-97	339.7
25	Feb-97	440.4
26	Mar-97	315.9
27	Apr-97	439.3
28	May-97	401.3
29	Jun-97	437.4
30	Jul-97	575.5
31	Aug-97	407.6
32	Sep-97	682.0
33	Oct-97	475.3
34	Nov-97	581.3
35	Dec-97	646.9

```
In [47]: # data types
```

```
df_shampoosales.dtypes
```

```
Out[47]: Month          object
Shampoo Sales    float64
dtype: object
```


```
In [48]: df_shampoosales.dtypes
```

```
Out[48]: Month          object
Shampoo Sales    float64
dtype: object
```

## Part a

Create a time series plot of the data. Label the axes.

```
In [49]: df_shampoosales.set_index('Month', inplace=True)
plt.rcParams["figure.figsize"] = (15,5)
df_shampoosales.plot()
plt.ylabel('Shampoo sales')
plt.xticks(rotation = 45)
plt.show()
```



## Part b

Which of the four components (level, trend, seasonality, noise) are present in this series?

Components appearing in the shampoo sales time series:

1. Level
2. Trend
3. Noise

## Question 4

(6 pts) The file, Beverages\_Shipment\_2020.csv, contains the US beverage product shipments data.

### Importing, examining and cleaning the shampoo sales dataset

```
In [15]: # Import beverages dataset
```

```
df_beverages_shipment = pd.read_csv('Beverages_Shipment_2020.csv')  
df_beverages_shipment
```

```
Out[15]:
```

	Month	Dollars (in Millions)
0	Jan-92	3519
1	Feb-92	3803
2	Mar-92	4332
3	Apr-92	4251
4	May-92	4661
...	...	...
175	Aug-06	7039
176	Sep-06	6440
177	Oct-06	6446
178	Nov-06	6717
179	Dec-06	6320

180 rows × 2 columns


```
In [16]: # data types
```

```
df_beverages_shipment.dtypes
```

```
Out[16]: Month                object  
Dollars (in Millions)      int64  
dtype: object
```

```
In [17]: # Create a time series plot
```

```
df_beverages_shipment.set_index('Month', inplace=True)  
plt.rcParams["figure.figsize"] = (15,5)  
df_beverages_shipment.plot()  
plt.xticks(rotation = 45)  
plt.show()
```



## Part a


Is there seasonality in this time series?

Yes, there is seasonality in the time series.

## Part b

Find the sample autocorrelation function for this data set. (2 pts) (For Python, you can use the “plot\_acf” function in “statmodels” module. Plot at least 25-30 lags)

```
In [19]: sm.graphics.tsa.plot_acf(df_beverages_shipment['Dollars (in Millions)'].squeeze(), lags=40)
plt.show()
```



## Part c

From the autocorrelation plot in (b), what is the seasonal period?

The seasonal period is 12 months ie it is an annual seasonal trend. It can be observed from the ACF plot that the high peaks are at 12, 24, and so on which indicates a high correlation between the values at these lags.



## Question 5

(10 pts) Data on US coal production is given in Coal\_Production\_US\_2020.csv.

**Importing, examining and cleaning the coal production dataset**

```
In [22]: # Import the coal production dataset

df_coal_production = pd.read_csv('Coal_Production_US_2020.csv')
df_coal_production
```

Out[22]:

	Year	Coal Production, Short Tons in Thousands
0	1949	480570
1	1950	560388
2	1951	576335
3	1952	507424
4	1953	488239
5	1954	420789
6	1955	490838
7	1956	529774
8	1957	518042
9	1958	431617
10	1959	432677
11	1960	434329
12	1961	420423
13	1962	439043
14	1963	477195
15	1964	504182
16	1965	526954
17	1966	546822
18	1967	564882
19	1968	556706
20	1969	570978
21	1970	612661

<b>Year</b>	<b>Coal Production, Short Tons in Thousands</b>
<b>22</b> 1971	560919
<b>23</b> 1972	602492
<b>24</b> 1973	598568
<b>25</b> 1974	610023
<b>26</b> 1975	654641
<b>27</b> 1976	684913
<b>28</b> 1977	697205
<b>29</b> 1978	670164
<b>30</b> 1979	781134
<b>31</b> 1980	829700
<b>32</b> 1981	823775
<b>33</b> 1982	838112
<b>34</b> 1983	782091
<b>35</b> 1984	895921
<b>36</b> 1985	883638
<b>37</b> 1986	890315
<b>38</b> 1987	918762
<b>39</b> 1988	950265
<b>40</b> 1989	980729
<b>41</b> 1990	1029076
<b>42</b> 1991	995984
<b>43</b> 1992	997545
<b>44</b> 1993	945424
<b>45</b> 1994	1033504
<b>46</b> 1995	1032974
<b>47</b> 1996	1063856

	Year	Coal Production, Short Tons in Thousands
48	1997	1089932
49	1998	1117535
50	1999	1100431
51	2000	1073612
52	2001	1127689
53	2002	1094283
54	2003	1071753
55	2004	1112099
56	2005	1133253

```
In [23]: # data types
```

```
df_coal_production.dtypes
```


```
Out[23]: Year                                int64
Coal Production, Short Tons in Thousands    int64
dtype: object
```

## Part a

Plot the coal production data and the sample autocorrelation function.


```
In [25]: # Create a time series plot
```

```
df_coal_production.set_index('Year', inplace=True)  
plt.rcParams["figure.figsize"] = (15,5)  
df_coal_production.plot()  
plt.xticks(rotation = 45)  
plt.show()
```



```
In [27]: # autocorrealtion function plot
```

```
sm.graphics.tsa.plot_acf(df_coal_production['Coal Production, Short Tons in Thousands'].squeeze(), lags=40, plt.show())
```



## Part b

Is the time series stationary or non-stationary ?


It can be observed from the time series plot that the time series is non-stationary due to the fact that it does not oscillate around a constant mean. However, it can be ascertained further from the ACF plot. There is a gradual decay in the values and the correlation tapers off slowly as the lags increase. This indicates that the time series is non-stationary.

## Part c


Plot the first difference of the time series and the sample autocorrelation function of the first difference.

```
In [33]: # plotting 1st order difference
```

```
ts_diff = df_coal_production.diff()  
ts_diff.plot()  
plt.show()
```



```
In [38]: # autocorrealtion function plot of first difference  
  
sm.graphics.tsa.plot_acf(ts_diff[1:].squeeze(), lags=40)  
plt.show()
```





## Part d

What impact has differencing had on the time series? Comment with respect to presence or absence of stationarity

There is stationarity in the time series after differencing. This is evident from the sharp drops and peaks in the ACF plot instead of a gradual decay.